

안녕하세요!

CS231n 딥러닝 스터디

1주차_1차시

| 201711703 김지우 |

INDEX



Introduction



Image Classification Pipeline





Loss Functions



Computer Vision




Visual Data



Algorithm

History of Computer Vision

- ⊙ Block World
- ⊙ Stages of Visual Representation
- ⊙ Generalized Cylinder
- ⊙ Pictorial Structure




The object is composed of “simple geometric primitives”



History of Computer Vision


- ◎ Normalized Cut
- ◎ Face Detection



“Object Segmentation”

History of Computer Vision

- © SIFT
- © Spatial Pyramid Matching
- © Histogram of Gradients
- © Deformable Part Model



“Feature Based Object Recognition”


History of Computer Vision

- © PASCAL Visual Object Challenge
- © IMAGENET



“Real Object Recognition”

어떤 걸 배울까요?




“Convolutional Neural Networks(CNN)”

“Visual Recognition”,
그중에서도 “Image Classification”

Semantic Gap

The Problem: Semantic Gap



This image by Nikita is
licensed under [CC-BY 2.0](#)

```
[[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]  
[ 91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]  
[ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]  
[ 99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]  
[106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]  
[114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]  
[133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]  
[128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]  
[125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]  
[127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]  
[115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]  
[ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]  
[ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]  
[ 62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]  
[ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]  
[ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]  
[118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]  
[164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]  
[157 170 157 120 93 86 114 132 112 97 69 55 70 82 99 94]  
[130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]  
[128 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]  
[123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]  
[122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]  
[122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]
```

What the computer sees

An image is just a big grid of
numbers between $[0, 255]$:

e.g. $800 \times 600 \times 3$
(3 channels RGB)




Challenges

- ◎ Viewpoint Variation
- ◎ Illumination
- ◎ Deformation
- ◎ Occlusion
- ◎ Background Clutter
- ◎ Intraclass Variation



Data-Driven Approach

1. Collect a dataset of images and labels
2. Use machine learning to train a classifier
3. Evaluate the classifier on new images



“Classifier”가 필요하다!

Nearest Neighbor

How to?

1. Memorize all data and labels
2. Predict the label of the most similar training image

가장 비슷한 이미지를 어떻게 찾을까?

"Distance Metric"을 이용하고,
"L1 distance"를 사용한다.

그러나, 이 방법은 Classifier를 Training하는데 오래 걸리지 않지만,
Prediction하는데 오래 걸리기 때문에 **좋은 방법은 아니다.**

K-Nearest Neighbors

How to?

Instead of copying label from nearest neighbor, take majority vote from K closet points


장점?

K-Nearest Neighbors를 이용하면, decision boundary를 부드럽게 만들어주고 더 나은 결과로 이끌어준다.

Distance Metric

L1 (Manhattan) distance


$$d_1(I_1, I_2) = \sum_p |I_1^p - I_2^p|$$



1. 좌표계 선택에 따라 달라진다.
2. 좌표축을 회전시키면, 점들 사이의 거리가 달라진다.
3. 벡터의 각각의 entry들이 중요한 의미를 가질 때 사용한다.

L2 (Euclidean) distance


$$d_2(I_1, I_2) = \sqrt{\sum_p (I_1^p - I_2^p)^2}$$



좌표축을 바꿔도 L2 distance는 달라지지 않는다.




Distance Metric



Setting Hyperparameters


- ◎ Data를 **Train, Validation, Test** 데이터로 나눠 Training Data를 통해 알고리즘을 학습시키고, Validation Data를 통해 Hyperparameter를 선택하고, Test Data를 통해 알고리즘이 새로운 데이터에서 얼마나 잘 작동하는 지 확인한다.
- ◎ **Cross-Validation**은 Training Data를 n 개의 부분으로 나누어 $n-1$ 개를 Training Data로 활용하고, 나머지 1개를 Validation Data로 활용하여 n 번의 Validation을 할 수 있는 방법을 말한다. 그러나 주로 데이터셋의 크기가 작을 때 유용하기 때문에, 딥러닝에는 잘 쓰이지 않는다.





Never Used

1. Very slow at test time
2. Distance metrics on pixels are not informative
3. Curse of dimensionality



“K-Nearest Neighbor on images never used”



Parametric Approach

Data-Driven Approach

1. Collect a dataset of images and labels
2. Use machine learning to train a classifier
3. Evaluate the classifier on new images



Parametric Approach

1. Summarize our knowledge of the training data and stick all that knowledge into these parameters
2. At test time, we only need these parameters(=weights), W

“테스트 할 때, 더 이상 실제 트레이닝 데이터가 필요하지 않다.”

Linear Classifier

$$f(x, W) = Wx + b$$



X: input data(=training data)

W: parameters or weights

b: bias term(=a constant vector of elements that does not interact with training data)

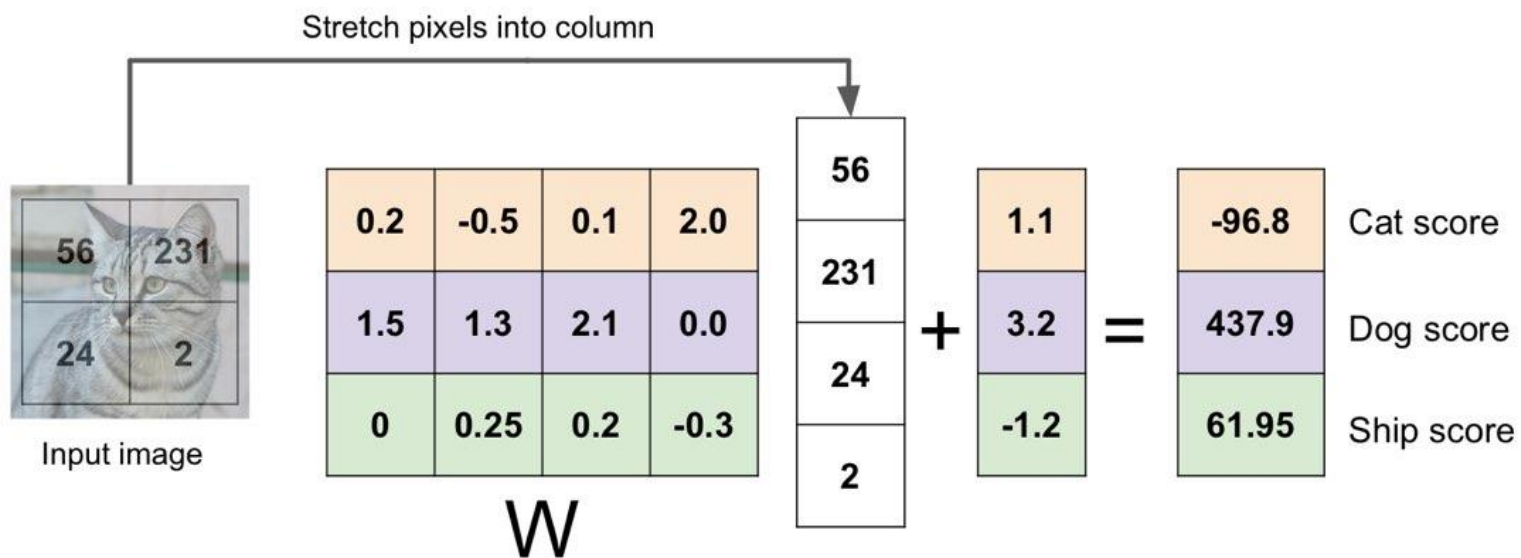


“Giving class scores”

“Data independent preferences for some classes over another”

Linear Classifier

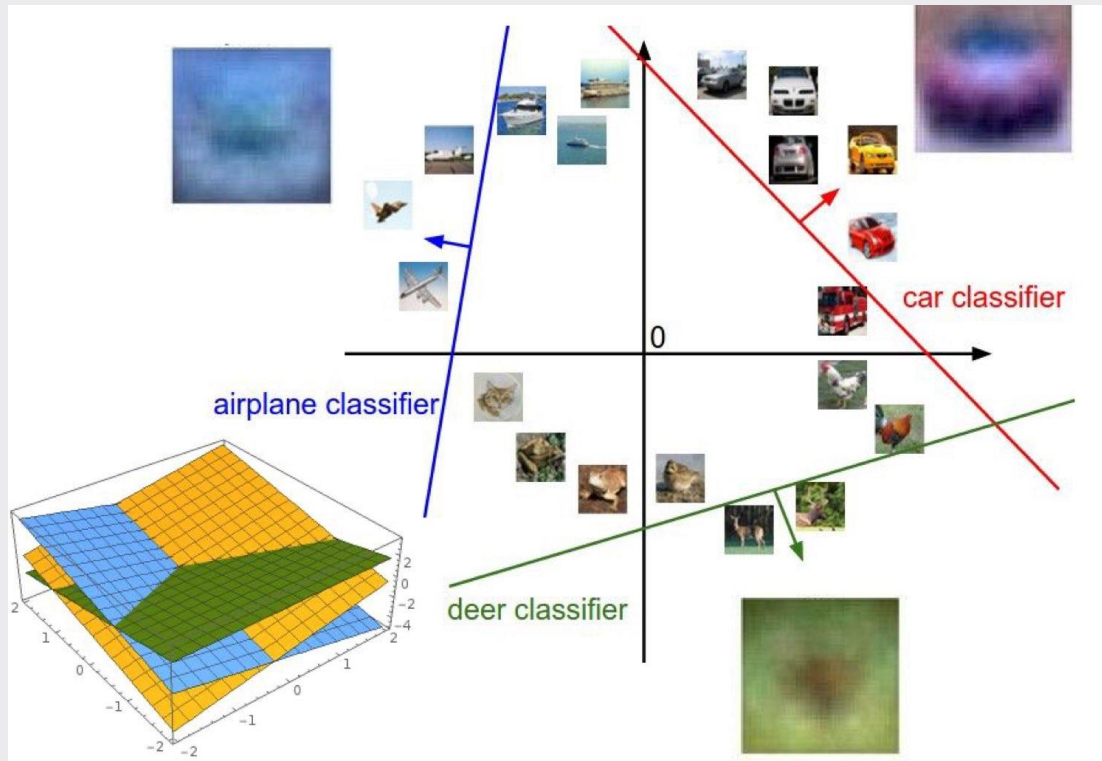
Example with an image with 4 pixels, and 3 classes (cat/dog/ship)




◎ W matrix의 각 행은 각 class의 template에 상응한다.

◎ W와 x를 내적한 것은, class(종류,분류)의 template과 image간의 유사성을 나타낸다.

Linear Classifier



“Draw linear separation between one category and the rest of categories”



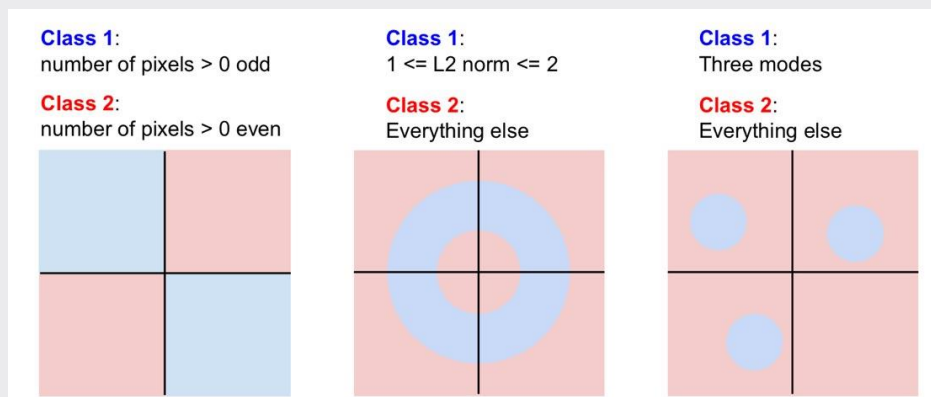
“linear decision boundaries”

Linear Classifier

Problem

The linear classifier is only learning **one template for each class**. So if there's sort of **variations** in how class might appear, it's trying to **average out** all of those different variations.

Hard cases for a linear classifier




"Parity problem"
"Multimodal data"



Loss function

“We need to quantify the badness of any particular W .”




“Loss Function”


Loss function

$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i)$$

“Loss over the dataset is a average of these losses summed over that entire data set.”



Hinge Loss (SVM Classifier)



Cross-Entropy Loss (Softmax Classifier)

Hinge Loss

$$\begin{aligned} L_i &= \sum_{j \neq y_i} \begin{cases} 0 & \text{if } s_{y_i} \geq s_j + 1 \\ s_j - s_{y_i} + 1 & \text{otherwise} \end{cases} \\ &= \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1) \end{aligned}$$

“ If the score for the correct category is greater than the score of the incorrect category by some safe margin(=1), if that’s the case that means the score for the true category is much larger than any of false categories, we’ll get a loss of zero.”

Hinge Loss

$$f(x, W) = Wx$$

$$L = \frac{1}{N} \sum_{i=1}^N \sum_{j \neq y_i} \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + 1)$$


"L=0으로 만드는 W는 unique하지 않다 "

Regularization

$$L(W) = \underbrace{\frac{1}{N} \sum_{i=1}^N L_i(f(x_i, W), y_i)}_{\text{Data loss}} + \underbrace{\lambda R(W)}_{\text{Regularization}}$$

Data loss: Model predictions should match training data

Regularization: Model should be “simple”, so it works on test data



In common use:

L2 regularization

$$R(W) = \sum_k \sum_l W_{k,l}^2$$

L1 regularization

$$R(W) = \sum_k \sum_l |W_{k,l}|$$

Elastic net (L1 + L2)

$$R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$$

Max norm regularization (might see later)

Dropout (will see later)

Fancier: Batch normalization, stochastic depth

Softmax Classifier

scores = unnormalized log probabilities of the classes.


$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}} \quad \text{where} \quad s = f(x_i; W)$$

Want to maximize the log likelihood, or (for a loss function) to minimize the negative log likelihood of the correct class:

$$L_i = -\log P(Y = y_i|X = x_i)$$

in summary:
$$L_i = -\log\left(\frac{e^{sy_i}}{\sum_j e^{s_j}}\right)$$

“Encourage our computed probability distribution that’s coming out of this softmax function to match this target probability distribution”



Softmax Classifier

Softmax Classifier (Multinomial Logistic Regression)



$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

unnormalized probabilities

cat
car
frog

3.2
5.1
-1.7

exp

24.5
164.0
0.18

normalize

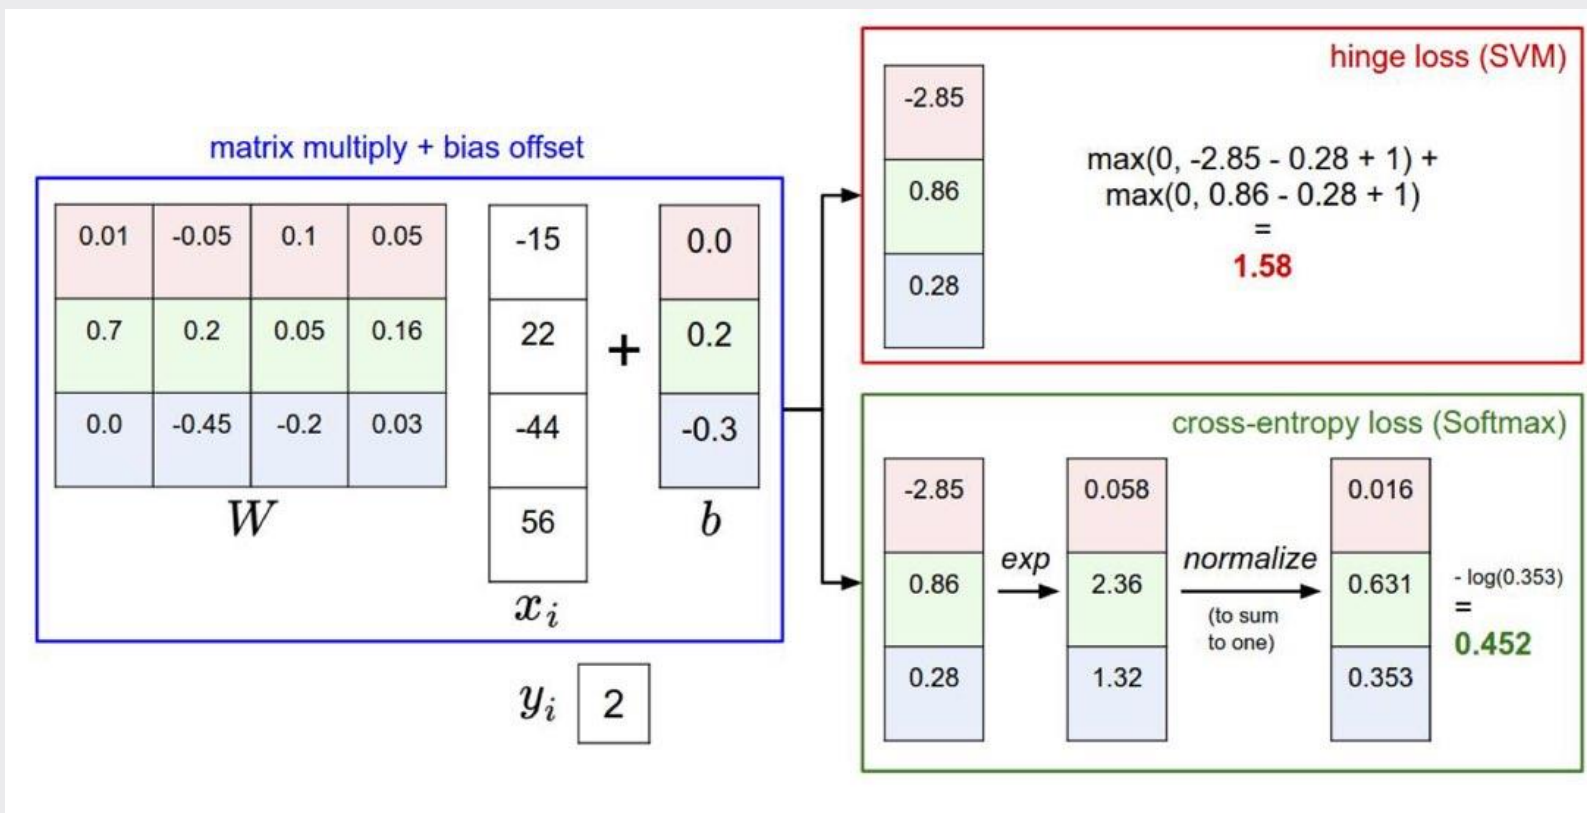
0.13
0.87
0.00

$L_i = -\log(0.13)$
 $= 0.89$

unnormalized log probabilities

probabilities

Difference between the two loss functions



“The difference between the two loss functions is how we choose to interpret those scores to quantitatively measure the badness afterwards.”

SVM: 바르게 분류되고 나면 더 이상 신경쓰지 않는다.

Softmax: 맞게 분류되더라도 맞는(올바른) class에 확률 질량이 더 많이 모여지는 방식으로 작동한다.

감 사 합 니 다 !

발표 들어주셔서 감사합니다
